

## Gravel2 – Task Scheduling and Concurrency

### 18-342: Fundamentals of Embedded Systems

**Rajeev Gandhi**

INI & ECE  
Carnegie Mellon University

**Carnegie Mellon**



#### Lab 4 Requirements

- Build a real-time operating system called Gravelv2
- Part I
  - Port your libc, swi and irq installation from lab 3
  - Create, schedule tasks according to rate monotonic policy
  - Create, manage mutexes
  - Schedule tasks according to a fixed pre-assigned priority
- Part II
  - Utilization bound (UB) test – to be covered in “Advanced Real-Time” lecture
  - Highest locker priority protocol – to be covered in “Advanced Real-Time” lecture

## Quick Look at Gravelv2

- Preemptive scheduling
- Maximum of 64 tasks
  - 63 “real” tasks, and one idle task
    - If you reserve the highest priority for special conditions (as in part 2), then 62 “real” tasks
  - Idle task executes when all other tasks are waiting on events
  - Idle task has the lowest task priority
  - Tasks are periodic (hence they never finish)
- Graavelv2 priorities
  - Higher numbers  $\Rightarrow$  lower priorities
  - Highest-priority task has priority 0
  - Two tasks should not have the same priority – you need to make sure of this

## Tasks

- A task in Gravelv2 is typically an infinite loop function

```
void MyTask(void *pdata)
{
    for(;;) {
        /* some application code */
        /* call one of Gravelv2' services: */

        mutex_lock(...);

        /* some application code */
        /* call one of Gravelv2' services: */

        mutex_unlock(...);

        /* more application code and Gravelv2 calls */
        event_wait(dev_num);
    }
}
```

## Task Creation

- Application may create tasks as follows

```
int main(int argc, char** argv)
{
    task_t tasks[2];
    tasks[0].lambda = fun1;
    tasks[0].data = (void*)'@';
    tasks[0].stack_pos = (void*)0xa2000000;
    tasks[0].C = 1;
    tasks[0].T = PERIOD_DEV0;

    tasks[1].lambda = fun2;
    tasks[1].data = (void*)'<';
    tasks[1].stack_pos = (void*)0xa1000000;
    tasks[1].C = 1;
    tasks[1].T = PERIOD_DEV1;

    task_create(tasks, 2);

    puts("Elvis could not leave the bus");
    return 0;
}
```

```
/* fun2 is defined as a function in the same file as
   main
*/
void fun2(void* str)
{
    while(1)
    {
        putchar((int)str);
        if (event_wait(1) < 0)
            panic("Dev 1 failed");
    }
}
```

## Task Creation

- Since tasks are periodic, they should never end
  - Hence no need for the `exit` syscall
- Once `task_create` function is called, you can assume that all the old task cease to exist (scheduler does not care about them anymore)
  - `task_create` function should never return
- Inside the `task_create` syscall, your code will look at the `task_t` data structure to learn everything needed about the tasks
- **Warning: Don't assume that user code will be sane**
  - **Your kernel should deal with unexpected input not conforming to specifications**

## What does dev data structure do in device.c

- Problem: How do we create periodic tasks? What do we do with a task after it has finished one of its instances?
    - Suspend the task until the next time period
    - Place the suspended task (it's TCB) in a sleep\_queue
    - Create events to occur periodically (with periodicity of the task)
    - On each event, check which tasks should be re-instantiated
  - dev data structure does precisely that
- ```
struct dev
{
    tcb_t* sleep_queue;
    unsigned long next_match;
};
```
- dev\_wait is called by tasks (through event\_wait) once they have finished execution of their current period and want to suspend themselves until the next period
  -

## What does dev data structure do in device.c

- dev\_update should be called by your timer interrupt handler code
  - dev\_update should check whether the next event for every device has occurred
  - If the next event has occurred
    - Wake up all the tasks on this device's sleep\_queue
    - Make these tasks ready to run

## SWI Handling

- You will port your SWI Handling code from Lab 3 to this lab
- Need to add support for additional SWIs as documented in Gravelv2 API
  - Need to remove `exit` syscall as well
- Difference from previous lab
  - Interrupts (IRQs) should be enabled when user programs make syscalls
  - Your kernel should be preemptible
- Modify your top level SWI handler to
  - Enable interrupts (at a safe point) before executing the actual syscall code
  - Disable interrupts (if needed, to prevent concurrency issues)

## IRQ Handling

- Since IRQs can occur while kernel is running, the IRQ handler is modified to run in SWI mode
  - After an IRQ occurs and context (not necessarily user context) has been saved, we switch to SVC mode
- The top-level interrupt handler has been provided to you (called `irq_wrapper` in file `int_asm.S`)
- You should understand what is going on in `irq_wrapper` and then modify your swi handler (top level) if needed
- Even though interrupts are handled in SWI mode, does not mean we allow nested interrupts

## Task Scheduling

- Each task in Gravelv2 is maintained using a Task Control Block (TCB)
  - Since the number of tasks in Gravelv2 is fixed, TCBs for all tasks are statically allocated as an array (`system_tcb`) in `sched.c`
  - The TCB structure is declared in `tasks.h`

```
struct tcb
{
    uint8_t      native_prio;
    uint8_t      cur_prio;
    sched_context_t context;
    int          holds_lock;
    volatile struct tcb* sleep_queue;
    /** Embed the kernel stack here -- AAPCS wants 8 byte alignment */
    uint32_t      kstack[OS_KSTACK_SIZE/sizeof(uint32_t)]
                  __attribute__((aligned(8)));
    uint32_t      kstack_high[0];
};
typedef volatile struct tcb tcb_t;
```

## Task Scheduling


- When `task_create` is executed, you must execute a function (`allocate_tasks` in `sched.c`)
  - In this function, you will
    - Set up the TCBs for all the tasks
    - Put the task in `runqueue` (make it available for running)
    - Set up the TCB for idle task
    - Make the idle task schedulable
- After all TCBs have been set up, you must context switch to the highest priority task that has been setup
  - This should occur in a function called `dispatch_nosave` in (`ctx_switch.c`)
  - All the dispatch functions are used for context switching in specific scenarios
    - `dispatch_save`: save the context of current task and context switch to highest priority task
    - `dispatch_nosave`: context switch to the highest priority task without worrying about saving current task's context
    - `dispatch_sleep`: save current task's context, make the current task not runnable and context switch

## Context Switching

- You will need to write code for context switching
- These functions are to be written in `ctx_switch_asm.S`
  - `ctx_switch_full` will get called from `dispatch_save` and `dispatch_sleep`
  - `ctx_switch_half` will get called from `dispatch_nosave`
- An assembly function `launch_task` has been provided to you
  - Understand this function
  - You will need to call this function the first time a task is launched
    - Don't define global/static variables called `first` or something similar
    - Once you understand what is going on in the code, it might help you with determining how context of a task will be initialized (before a task is launched for the first time)

## Run queues

- Gravelv2 maintains the list of currently runnable tasks on run queues
- Typically there will be a run queue consisting of tasks that are ready to run
  - On a context switch, the scheduler can go through this run queue and look for the highest priority task that is ready to run
- Maintain run queue as a linked list of TCBs of all tasks in your system



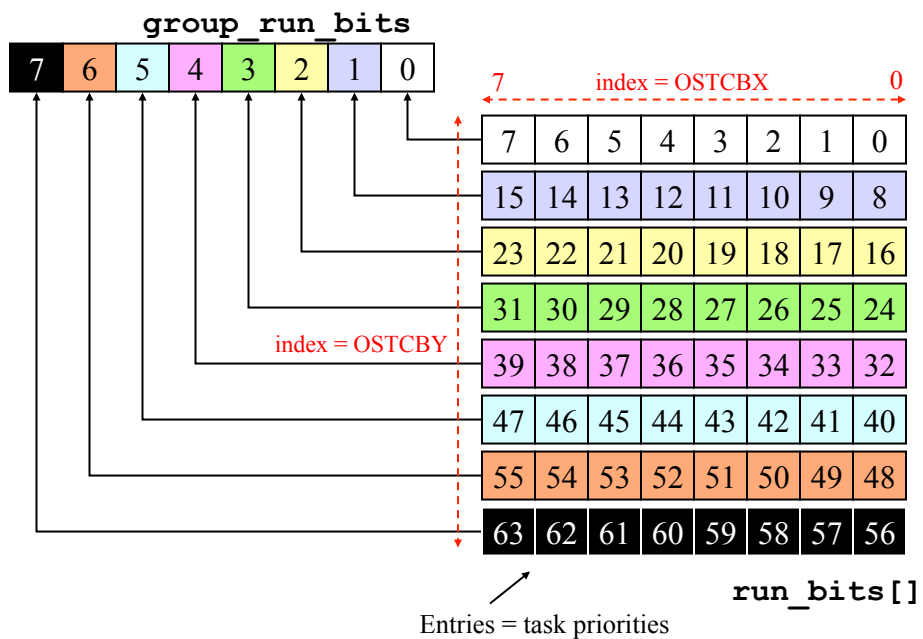
- How do you find the highest priority task in this list?
- Disadvantages?

## Runqueue in Gravelv2

Where is the runqueue maintained in Gravelv2?

Using the array `run_bits[8]` and variable `group_run_bits` in `run_queue.c`


## Ready Queue





## Task Groups

- Gravely2 divides its 64 priority levels into 8 groups
  - How do you determine a task's group?
    - priority >> 3**
    - This value is stored in the **OSTCBY** variable
      - Task priorities 0, 1, 2, ..., 7  $\Rightarrow$  task group = 0
      - Task priorities 8, 9, 10, ..., 15  $\Rightarrow$  task group = 1
    - i.e., **OSTCBY should** contain the 3 MSBs of the task's priority number
  - What is the task's position within the group?
    - priority & 0x07**
    - This value is stored in the **OSTCBX** field of the task's OS\_TCB
    - OSTCBX** contains the 3 LSBs of the task's priority number



$$\text{priority} = \text{OSTCBY} \ll 3 + \text{OSTCBX}$$

## Ready List

- Gravely2 maintains two variables `group_run_bits` and `run_list[]` which are used to store the **list of tasks ready to run**
- When a task is created or is made ready to run, the appropriate bits of `group_run_bits` and `run_list[]` are set to indicate to the scheduler that the task is ready to run
- Example: Assume `group_run_bits` and `run_list[8]={0,0,...0}` and a task of priority `prio = 17` is created, in order to let the scheduler know the presence of the task at `prio = 17`, `TaskCreate` must
  - Compute priority group (OSTCBY field) of the task
 
$$\text{OSTCBY} = (\text{prio} \gg 3) = 2$$
  - Compute task's position in the priority group (OSTCBX field)
 
$$\text{OSTCBX} = (\text{prio} \& 0x07) = 1$$
  - Set bit number OSTCBY of `group_run_bits` to 1
  - Set bit number OSTCBX of `run_bits[OSTCBY]` equal to 1

## Example

- What will be the values of `group_run_bits` and `run_bits[]` when 3 tasks of priorities 16, 17 and 25 are created?
- Assume that initially `group_run_bits = 0`, and `run_bits = {0, 0, 0, 0, 0, 0, 0, 0}`

## Finding the Highest-Priority Ready Task

- Suppose that you want to find the highest-priority task (let's denote its priority by `prio`) that is ready to run
- Can be done quickly (without scanning the entire linked list of created tasks) through `group_bits` and `run_bits[]`

– Algorithm:

- Find the least significant bit set in `group_bits` (call it `y`)
- Find the least significant bit set in `run_bits[y]` (say `x`)
- Highest priority task ready to run is  $(y \ll 3) + x$

- Optimization – Since you will be doing this often, to make this process faster, there is a special table called `prio_unmap_table[]` with 256 entries that helps you do the above

```
y = prio_unmap_table[group_bits];
x = prio_unmap_table[run_bits[y]]
prio = (y << 3) + x
```

- You are likely to need this kind of code in `run_queue()`

## Finding the Highest-Priority Ready Task (contd.)

- Assume group\_bits=12 and run\_bits={0,0,3,2,0,0,0,0}
- How do you determine the highest priority task to run

```
UnMapTbl[] = {
  7, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  6, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  7, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  6, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0,
  4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0
};

y = UnMapTbl[group_bits]
x = UnMapTbl[run_bits[y]]
prio = (y << 3) + x

Once prio is known, you
can get
to the TCB of this task by
accessing system_tcb[prio]
```

- Question: Do you need to do anything to group\_run\_bits and run\_bits[] when suspending a task?

## Summary

- Lab 4 overall summary
- Please follow course mailing list for more details
  - You may have several questions
  - Please use mailing list rather than sending individual requests
  - You may be helping others with similar problem
- Aim to finish task management first – you should try and use what we've covered today to do this as soon as possible
  - Lab is even more challenging than Lab 3
  - Very useful to understand embedded OS concepts if you have never done a course on OS