

## Week 3. ARM Assembly Programming

### 18-342: Fundamentals of Embedded Systems

**Rajeev Gandhi**

INI & ECE  
Carnegie Mellon University

**Carnegie Mellon**



### Overview of this Week

- Conditional execution of instructions
  - Improves code density
  - Can improve performance
- Load/Store instructions
  - Different addressing modes
- Multiple register load and store instructions
- Performing stack operations using load/store instructions
- Digging deeper
  - Limitations of branch instructions
  - Loading arbitrary 32-bit constants in registers
- **Announcement:**
  - Quiz on Thursday

## Value of Conditional Execution

- ◆ This improves code density *and* performance by reducing the number of forward branch instructions

```
if (x != 0)      CMP    r3, #0
    a = b+c;      BEQ    skip
else            ADD    r0, r1, r2
    a = b-c;      B      afterskip
                skip
                SUB    r0, r1, r2
                afterskip
```

```
CMP    r3, #0
ADDNE  r0, r1, r2
SUBEQ  r0, r1, r2
```

- ◆ Can also be used to optimize a countdown loop (where the loop variable decrements from a positive number to zero)

```
loop
    . . .
    SUB r1, r1, #1
    CMP r1, #0
    BNE loop
```

```
loop
    . . .
    SUBS r1, r1, #1
    BNE loop
```

3

## Examples of Conditional Execution

- ◆ Use a sequence of several conditional instructions

```
if (a==0) x=1;
    CMP    r0, #0
    MOVEQ   r1, #1
```

- ◆ Set the flags, then use various condition codes

```
if (a==0) x=0;
if (a>0)  x=1;
    CMP    r0, #0
    MOVEQ   r1, #0
    MOVGT   r1, #1
```

- ◆ Use conditional compare instructions

```
if (a==4 || a==10) x=0;
    CMP    r0, #4
    CMPNE   r0, #10
    MOVEQ   r1, #0
```

4

## Single Register Data Transfer

### ◆ ARM is based on a “load/store” architecture

- All operands should be in registers
- Load instructions are used to move data from memory into registers
- Store instructions are used to move data from registers to memory
- Flexible – allow transfer of a word or a half-word or a byte to and from memory

LDR/STR	Word
LDRB/STRB	Byte
LDRH/STRH	Halfword
LDRSB	Signed byte load
LDRSH	Signed halfword load

### ◆ Syntax:

- LDR{<cond>}{<size>} Rd, <address>
- STR{<cond>}{<size>} Rd, <address>

5

## LDR and STR

- ◆ LDR and STR instructions can load and store data on a boundary alignment that is the same as the datatype size being loaded or stored
- ◆ LDR can only load 32-bit words on a memory address that is a multiple of 4 bytes – 0, 4, 8, and so on

### ◆ LDR r0, [r1]

- Loads register r0 with the contents of the memory address pointed to by register r1

### ◆ STR r0, [r1]

- Stores the contents of register r0 to the memory address pointed to by register r1

- ◆ Register r1 is called the [base address register](#)

6

## Addressing Modes

- ◆ ARM provides three addressing modes
  - Preindex with writeback
  - Preindex
  - Postindex
- ◆ Preindex mode useful for accessing a single element in a data structure
- ◆ Postindex and preindex with writeback useful for traversing an array

7

## Addressing Modes

- ◆ Preindex
  - Same as preindex with writeback, but does not update the base register
  - Example: `LDR r0, [r1, #4]`
- ◆ Preindex with writeback
  - Calculates address from a base register *plus* address offset
  - Updates the address in the base register with the new address
  - The *updated base register value* is the address used to access memory
  - Example: `LDR r0, [r1, #4]!`
- ◆ Postindex
  - Only updates the base register *after* the address is used
  - Example: `LDR r0, [r1], #4`

8

## Examples

```
PRE   r0 = 0x00000000
      r1 = 0x00009000
      mem32[0x00009000] = 0x01010101
      mem32[0x00009004] = 0x02020202
```

### Preindexing with writeback

```
LDR r0, [r1, #4]!
```

```
POST  r0 =
      r1 =
```

### Preindexing

```
LDR r0, [r1, #4]
```

```
POST  r0 =
      r1 =
```

### Postindexing

```
LDR r0, [r1], #4
```

```
POST  r0 =
      r1 =
```

## More on Addressing Modes

- ◆ Address <address> accessed by LDR/STR is specified by
  - A base register plus an offset
- ◆ Offset takes one of the three formats
  1. **Immediate**: offset is a number that can be added to or subtracted from the base register
 

Example: LDR r0, [r1, #8];                    r0 ↪ mem[r1+8]  
           LDR r0, [r1, #-8];                  r0 ↪ mem[r1-8]
  2. **Register**: offset is a general-purpose register that can be added to or subtracted from the base register
 

Example: LDR r0, [r1, r2];                    r0 ↪ mem[r1+r2]  
           LDR r0, [r1, -r2];                  r0 ↪ mem[r1-r2]
  3. **Scaled Register**: offset is a general-purpose register shifted by an immediate value and then added to or subtracted from the base register
 

Example: LDR r0, [r1, r2, LSL #2];            r0 ↪ mem[r1+4\*r2]

## Multiple-Register Transfer

- ◆ Load-store-multiple instructions can transfer multiple registers between memory and the processor in a single instruction
- ◆ Advantages
  - More efficient than single-register transfers for moving blocks of data around memory
  - More efficient for saving and restoring **context** and **stacks**
- ◆ Disadvantages
  - ARM does not interrupt instructions when executing ⇒ load-store multiple instructions can increase interrupt latency
- ◆ Compilers can limit interrupt latency by providing a switch to control the max number of registers that can be transferred on a load-store-multiple

```
LDM<cond><addrMode> Rn{!}, <registerList>{^}  
STM<cond><addrMode> Rn{!}, <registerList>{^}
```

11

## More on Load-Store-Multiple

- ◆ Transfer occurs from a base-address register  $R_n$  pointing into memory
- ◆ Transferred registers can be either
  - Any subset of the current bank of registers (default)
  - Any subset of the user mode bank of registers when in a privileged mode (postfix instruction with a '^')
  - Processor not in user mode or system mode
  - Writeback is not possible, i.e., ! cannot be supported at the same time
  - If  $pc$  is in the list of registers, additionally copy  $spsr$  to  $cpsr$
- ◆ Register  $R_n$  can be optionally updated following the transfer
  - If register  $R_n$  is followed by the ! character
- ◆ Registers can be individually listed or lumped together as a range
  - Use a comma with “{” and “}” parentheses to list individual registers
  - Use a “-” to indicate a range of registers
  - Good practice to list the registers in the order of increasing register number (since this is the usual order of memory transfer)

12

## Addressing Modes for Load-Store-Multiple

- ◆ Suppose that  $N$  is the number of registers in the list of registers
- ◆ IA (increment after)
  - Start reading at address  $R_n$ ; ending address is  $R_n + 4N - 4$
  - $R_n!$  equals  $R_n + 4N$
- ◆ IB (increment before)
  - Start reading at address  $R_n + 4$ ; ending address is  $R_n + 4N$
  - $R_n!$  equals  $R_n + 4N$
- ◆ DA (decrement after)
  - Start reading at address  $R_n - 4N + 4$ ; ending address is  $R_n$
  - $R_n!$  equals  $R_n - 4N$
- ◆ DB (decrement before)
  - Start reading at address  $R_n - 4N$ ; ending address is  $R_n - 4$
  - $R_n!$  equals  $R_n - 4N$
- ◆ ARM convention: DB and DA are like loading the register list backwards from sequentially descending memory addresses

13

## Things to Remember

- ◆ Any register can be used as the base register
- ◆ Any register can be in the register list
- ◆ Order of registers in the list does not matter
- ◆ The lowest register always uses the lowest memory address *regardless of the order in which registers are listed in the instruction*
- ◆ LDM and STM instructions **only transfer words**
  - Unlike LDR/STR instructions, they don't transfer bytes or half-words
- ◆ Can specify range instead of individual registers
  - Example: `LDMIA r10!, {r12, r2-r7}`
- ◆ If the base register is updated (using `!`) in the instruction, then it cannot be a part of the register set
  - Example: `LDMIA r10!, {r0, r1, r4, r10}` is not allowed

14

## Examples

<pre> PRE    r0 = 0x00080010         r1 = 0x00000000         r2 = 0x00000000         r3 = 0x00000000         mem32[0x8001c] = 0x04         mem32[0x80018] = 0x03         mem32[0x80014] = 0x02         mem32[0x80010] = 0x01 </pre>	<p>r0 (original) →</p>	<table border="1"> <tr><td>0x00080020</td><td>0x05</td></tr> <tr><td>0x0008001c</td><td>0x04</td></tr> <tr><td>0x00080018</td><td>0x03</td></tr> <tr><td>0x00080014</td><td>0x02</td></tr> <tr><td>0x00080010</td><td>0x01</td></tr> <tr><td>0x0008000c</td><td>0x00</td></tr> </table>	0x00080020	0x05	0x0008001c	0x04	0x00080018	0x03	0x00080014	0x02	0x00080010	0x01	0x0008000c	0x00	
0x00080020	0x05														
0x0008001c	0x04														
0x00080018	0x03														
0x00080014	0x02														
0x00080010	0x01														
0x0008000c	0x00														

<pre> LDMIA r0!, {r1-r3}  POST   r0 =         r1 =         r2 =         r3 = </pre>	<pre> LDMIB r0!, {r1-r3}  POST   r0 =         r1 =         r2 =         r3 = </pre>
---	---

15

## Example 1: Saving & Restoring Registers

- ◆ Here's what we want to accomplish
- Save the contents of registers r1, r2 and r3 to memory
  - Mess with the contents of registers r1, r2 and r3
  - Restore the original contents of r1, r2 and r3 from memory & restore r0

<pre> PRE    r0 = 0x00009000         r1 = 0x09         r2 = 0x08         r3 = 0x07  ; store contents to memory STMIB r0!, {r1-r3} ; mess with registers r1, r2, r3 MOV r1, #1 MOV r2, #2 MOV r3, #3 ; restore original r1, r2, r3 LDMDA r0!, {r1-r3} </pre>	<p>r0 (original) →</p>	<table border="1"> <tr><td>0x0000900c</td><td></td></tr> <tr><td>0x00009008</td><td></td></tr> <tr><td>0x00009004</td><td></td></tr> <tr><td>0x00009000</td><td></td></tr> </table>	0x0000900c		0x00009008		0x00009004		0x00009000		
0x0000900c											
0x00009008											
0x00009004											
0x00009000											

ARM convention: Highest memory location maps to highest numbered register

16



## Example 1: Block Copying

- ◆ Here's what we want to accomplish
  - Copy blocks of 32 bytes from a source address to a destination address
  - r9 points to the start of the source data
  - r10 points to the start of the destination data
  - r11 points to the end of the source data

```
loop
    ; load 32 bytes from source address and update r9 pointer
    LDMIA r9!, {r0-r7}
    ; store 32 bytes to destination address and update r10 pointer
    STMIA r10!, {r0-r7}
    ; check if we are done with the entire block copy
    CMP r9, r11
    ; continue until done
    BNE loop
```

17

## Stack Operations

- ◆ ARM uses load-store-multiple instructions to accomplish stack operations
- ◆ Pop (removing data from a stack) uses load-multiple
- ◆ Push (placing data on a stack) uses store-multiple
- ◆ Stacks are ascending or descending
  - Ascending (A): Grow towards higher memory addresses
  - Descending (D): Grow towards lower memory addresses
- ◆ Stacks can be full or empty
  - Full (F): Stack pointer sp points to the last used or full location
  - Empty (E): Stack pointer sp points to the first unused or empty location
- ◆ Four possible variants
  - Full ascending (FA) – LDMFA & STMFA
  - Full descending (FD) – LDMFD & STMFD
  - Empty ascending (EA) – LDMEA & STMEA
  - Empty descending (ED) – LDMED & STMED

18

## Stacks on the ARM

- ◆ ARM has an ARM-Thumb Procedure Call Standard (ATPCS)
  - Specifies how routines are called and how registers are allocated
- ◆ Stacks according to ATPCS
  - Full descending
- ◆ What does this mean for you?
  - Use STMFD to store registers on stack at procedure entry
  - Use LDMFD to restore registers from stack at procedure exit
- ◆ What do these handy aliases actually represent?
  - STMFD = STMDB (store-multiple-decrement-before)
  - LDMFD = LDMIA (load-multiple-increment-after)

19

## Example

```
PRE    r1 = 0x00000002
        r4 = 0x00000003
        sp = 0x00080014
```

```
STMFD  sp!, {r1, r4}
```

sp  
(original)

0x00080018	0x05
0x00080014	0x04
0x00080010	Empty
0x0008000c	Empty

sp  
(final)

0x00080018	0x05
0x00080014	0x04
0x00080010	0x03
0x0008000c	0x02

20


## Stack Checking

- ◆ Three stack attributes to be preserved
- ◆ Stack base
  - Starting address of the stack in memory
  - If `sp` goes past the stack base, stack underflow error occurs
- ◆ Stack pointer (`sp`)
  - Initially points to the stack base
  - As data is inserted when a program executes, `sp` descends memory and points to top of the stack
- ◆ Stack limit (`sl`)
  - If `sp` passes the stack limit, a stack overflow error occurs
  - ATPCS: `r10` is defined as `sl`
    - If `sp` is less than `r10` after items are pushed on the stack, stack overflow occurs

21

## Digging Deeper

### ◆ Branch instructions



- When executing the instruction, the processor
  - Shifts the offset left two bits, sign extends it to 32 bits, and adds it to the `pc`
- This gives a 26-bit offset from the current instruction
- What is the maximum distance (in bytes) that we can jump to, in a branch?

22


## LDR Instruction

◆ `LDR rd, [rn, #offset]`      $rd \leftarrow \text{mem}[rn+\text{offset}]$

- Decimal numbers prefixed by #

◆ `rd, rn` can be any register (`r0 – r15`)

◆ Binary encoding of LDR/STR instruction (with immediate offset addressing mode)



- distinguish between load (1) and store (0)

- add or subtract offset

23

## LDR (contd.)

◆ LDR can be used for branches beyond the range of BL instruction

- Example: `LDR pc, [pc, #offset]`      $pc \leftarrow \text{mem}[pc+\text{offset}]$
- The address of the branch should be stored in memory location `curraddr + offset + 8`

◆ Example

```

0x10000000    add r0, r1, r2
0x10000004    ldr pc, [pc, #4];
0x10000008    sub r1, r2, r3
0x1000000c    cmp r0, r1
0x10000010    0x20000000
...
...
Branch_target
0x20000000    str r5, [r13, #-4]!
...
...

```

◆ STR/LDR instructions can be used to save/restore registers on function entry/exit

24

## Digging Deeper: Loading Constants

- ◆ There is no single instruction which will load a 32 bit immediate constant into a register without performing a data load from memory.
  - All ARM instructions are 32 bits long
  - ARM instructions do not use the instruction stream as data
- ◆ The data processing instruction format has 12 bits available for operand2
  - If used directly, this would only give a range of 4096 bytes
- ◆ Instead it is used to store 8-bit constants, giving a range of 0 - 255
- ◆ These 8 bits can then be rotated right through an even number of positions
  - This gives a much larger range of constants that can be directly loaded, though some constants will still need to be loaded from memory

25

## Specific Example: MOV

- ◆ MOV or MVN data processing instruction can be used to load 8-bit numbers into registers
  - `MOV r0, #0x07 ; r0=0x00000007`
- ◆ Use MOV with the barrel shifter to load more than 8-bit numbers into registers
  - `MOV r0, #0x0F, #2 ; r0=0xC0000003`
- ◆ Remember operand2 in MOV instruction takes 12 bits
  - 8 bits are used for the immediate value
  - 4 bit rotate value (which should be an even number between 0 and 30)
- ◆ Rule to remember is “8-bits rotated right by an even number of bit positions”.

26

## Loading Constants: Pseudo-Instruction LDR

- ◆ To allow larger constants to be loaded, the assembler offers a **pseudo-instruction**:
  - `LDR Rd, =const`
- ◆ This will either:
  - Produce a `MOV` or `MVN` instruction to generate the value (if possible)**or**
  - Generate a `LDR` instruction with a `pc`-relative address to read the constant from a *literal pool* (constant data area embedded in the code)
  - What is a literal pool?
    - Portion of memory used by the assembly code to store constants
- ◆ This is the recommended way of loading constants into a register

27

## Loading Constants – Example

### Original assembly code

```
LDR r0, =0x55555555
LDR r1, =0x4867
ADD r2, r1, r0
SWI 0x123456
```

### Machine code generated by assembler

```
0x00000000: e59f0008
0x00000004: e59f1008
0x00000008: e0812000
0x0000000c: ef123456
0x00000010: 55555555
0x00000014: 00004867
```

- Assembler places the constant `0x55555555` at location `0x00000010`
- and converts the pseudo-instruction to `LDR r0, [pc, #8]`
- Assembler places the constant `0x4867` at location `0x00000014`
- and converts the pseudo-instruction to `LDR r0, [pc, #8]`

28