

Week 14 – RTOS: Memory Management

18-342: Fundamentals of Embedded Systems

Rajeev Gandhi

INI & ECE
Carnegie Mellon University

Carnegie Mellon



Today's Lecture

- Memory Protection using the MPU
- Introduction to virtual memory
- How does virtual memory work?
- Single-level and two-level page tables
- Translation lookaside buffers
- A quick look at ARM's memory management unit

Memory Protection

- ARM processors provide two different ways of protecting a system i.e. by providing restricted access to system's resources (memory and peripheral devices)
 - MPU -- Memory Protection Unit
 - MMU -- Memory Management Unit
- Some ARM processors only have an MPU while others only have an MMU
- ARM's Memory Protection Unit
 - Allows physical memory to be divided into 8 or fewer sections or regions
 - Can specify the starting address and size of each section
 - Maximum Size -- 4GB
 - Minimum Size -- 4KB
 - Access permissions can be specified for each section
 - Regions can have privileged mode only access, privileged mode full access and user mode read-only access, full access


Regions



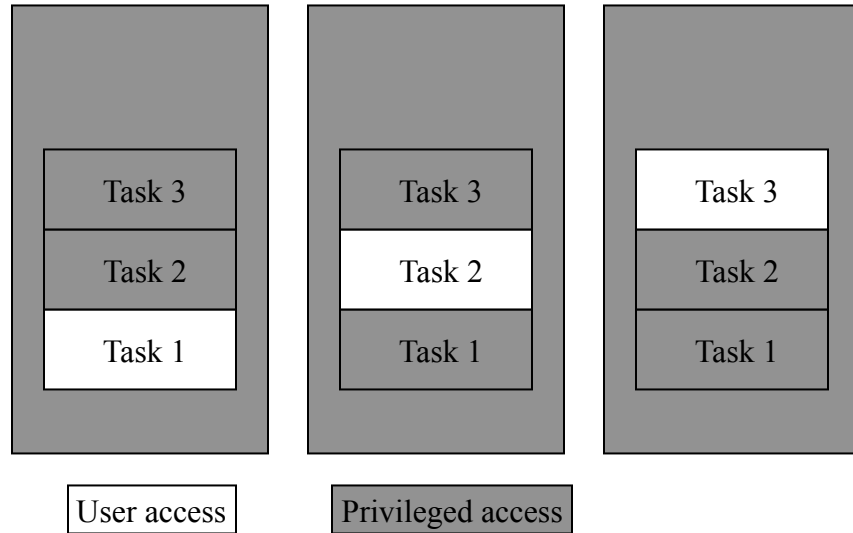
MPU: Regions

- A maximum of 8 regions can be defined
- Rules governing regions
 - Regions assigned privileges
 - Regions can overlap with other regions
 - Regions are assigned priorities independent of privilege level of the region
 - When regions overlap, the attributes of region with highest priority take precedence in the overlapped memory locations
 - A region's starting address must be a multiple of its size
 - Size of a region can be a power of two from 4KB to 4GB
 - Data or prefetch abort is generated when the processor accesses a region in a mode where it does not have access to a region

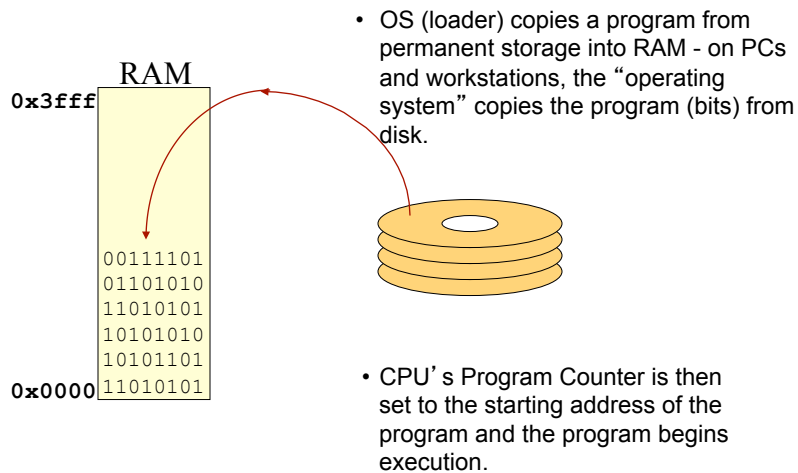
Overlapping Regions



Overlapping Regions




Virtual Memory Introduction – How Does a Program Start Running?



What if the program is too big?

- Some machines won't let you run the program
 - Original DOS
 - Why this limitation?

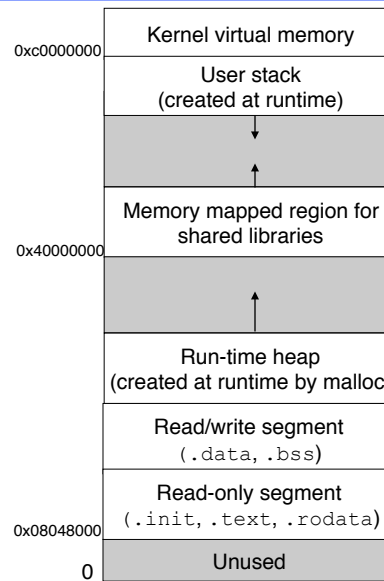


Solution: Virtual Memory

- What is **virtual memory**?
 - Technique that allows execution of a program that may not completely reside in memory (RAM)
- Allows the computer to “fake” a program into believing that its memory space is larger than physical RAM
- Key concept -- use DRAM as a cache for the disk
 - Address space of a process can exceed main memory size
 - Sum of address space of multiple processes can exceed main memory size
- Why is VM important?
 - Cheap - no longer have to buy lots of RAM
 - Removes burden of memory resource management from the programmer

But Wait There is More...


- Virtual Memory concept also
 - Simplifies memory management by providing each process with same uniform address space
 - Protects one process from corrupting data of another process or corrupting its own read only (text) sections
- In embedded systems virtual memory is typically used to provide memory protection and uniform address space for processes



Source: Bryant & O'Hallaron


Physical versus Virtual Addressing

- Two memory “spaces”
 - **Virtual memory space** - what the program “sees”
 - **Physical memory space** - what the program runs in (size of RAM)
- Virtual memory requires
 - Dedicated hardware on CPU chip called Memory Management Unit (MMU)
 - Cooperation between CPU hardware & operating system

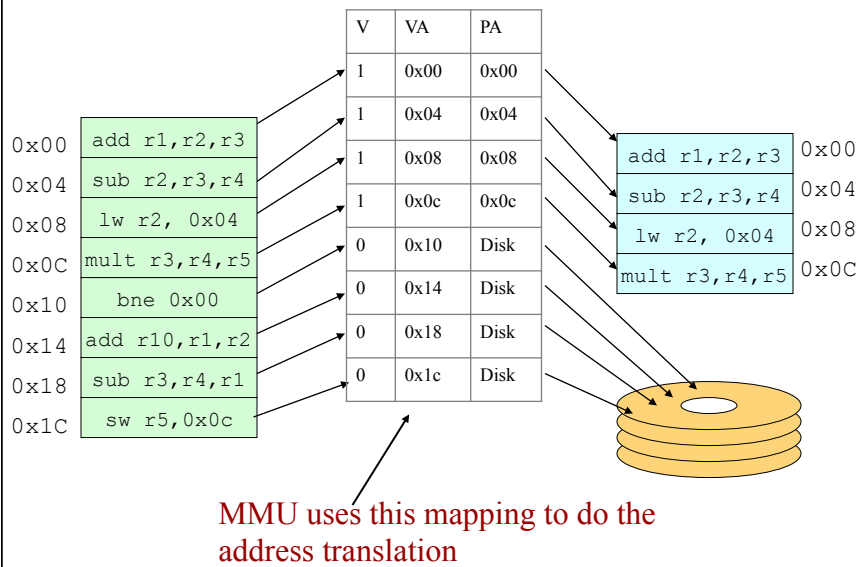


Source: Bryant & O'Hallaron

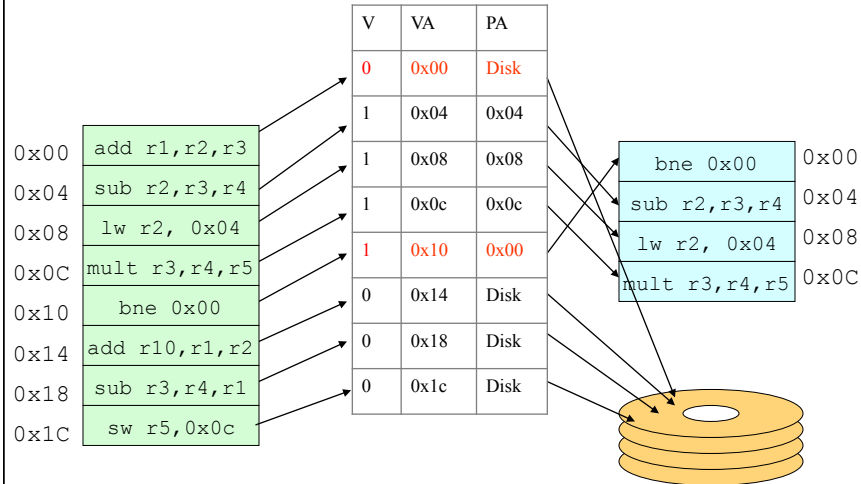
Example: Virtual and Physical Address Spaces



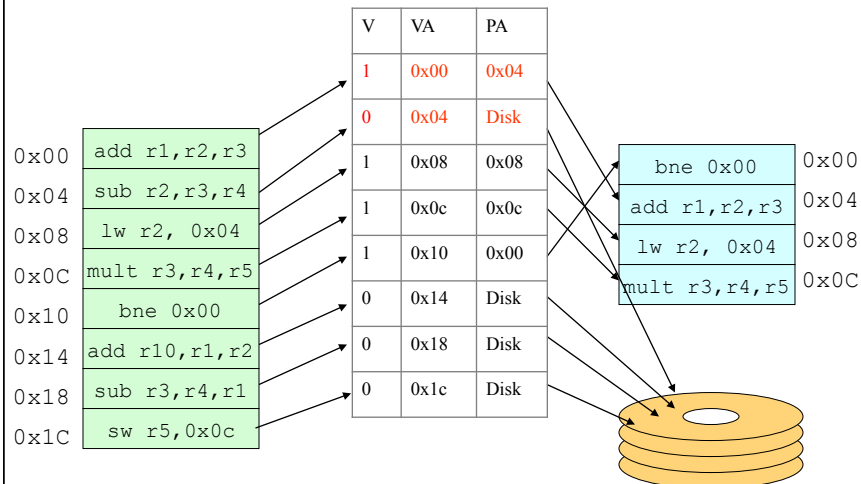
Example (con'td): Need VA-to-PA mappings



Example (cont' d): After handling a page fault

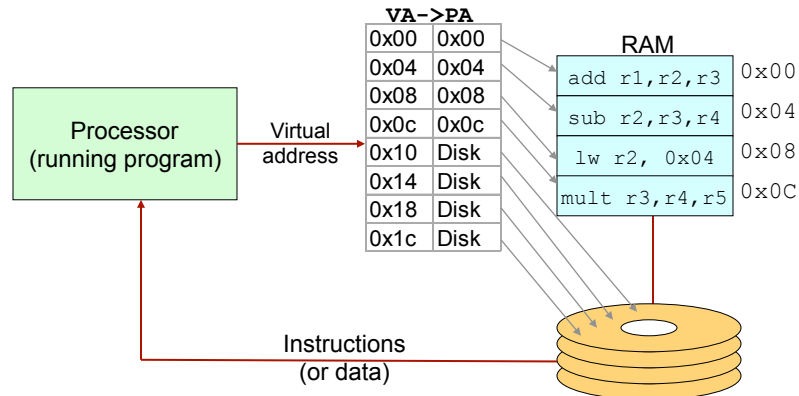


Example (con'td): After a second page fault



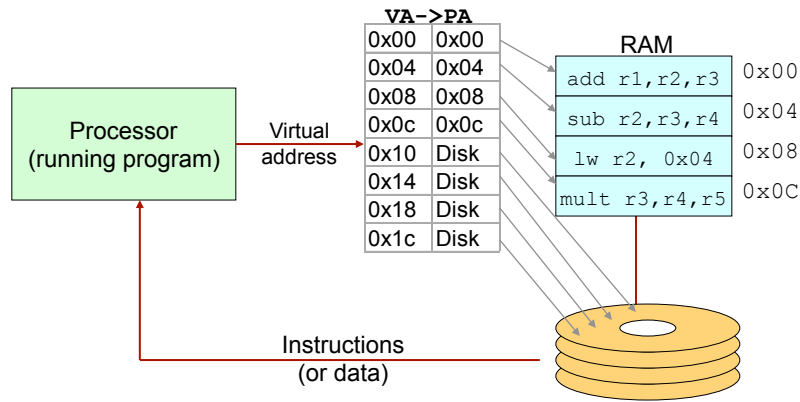
Basic VM Algorithm

- Program asks for virtual address
- MMU translates **virtual address** (VA) to **physical address** (PA)
- Computer reads PA from RAM, returning it to program



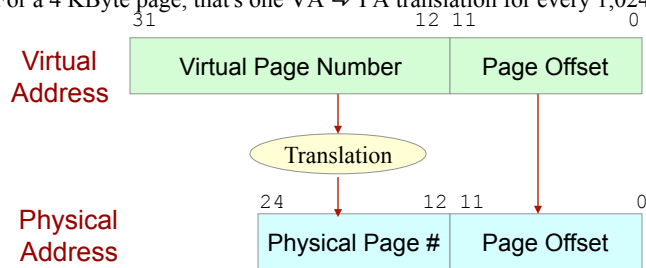
Page Tables

- Table which holds VA \Rightarrow PA translations is called the **page table**
- In our current scheme, each word is translated from a virtual address to a physical address
 - Where is the page table located?
 - How big is the page table (for our scheme assuming 32-bit addresses)?

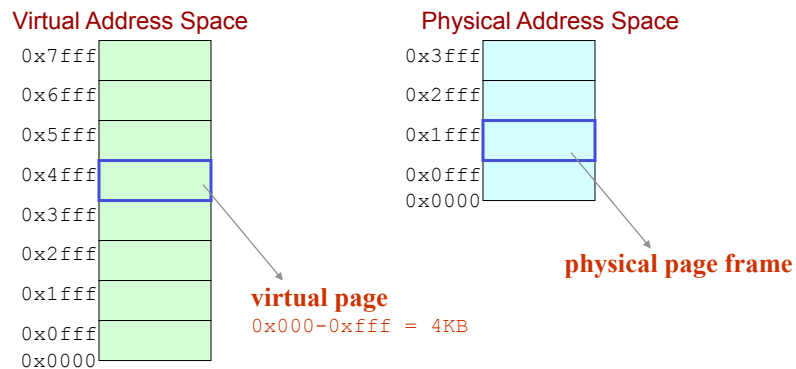


Real Page Tables

- Instead of the fine-grained VM where any virtual word can map to any RAM word location, partition memory into chunks called **pages**
 - Typical page sizes are 1, 4, 8, 64 Kbytes or 1 Mbyte
 - With 32-bit addresses and assuming 4 KByte page size, the 20 MSBs determine page number
- Map each virtual page to a physical page
 - **Within** a page, the virtual offset == physical offset
- This reduces the number of VA \Rightarrow PA translation entries
 - Only one translation per page
 - For a 4 KByte page, that's one VA \Rightarrow PA translation for every 1,024 words



Virtual Pages & Physical Page Frames



Every address within a virtual page maps to the same location within a physical page frame


- In other words, bottom $\log_2(\text{page size in bytes})$ bits are **not** translated

Page Table

- Who fills/sets up and manages the page table?
- How does the MMU know where the page table is located in memory?
 - A register (typically called page-table base register (PTBR)) is used to store the starting address of the table
- Virtual page number can be used as an index into this table to determine the corresponding physical page number
- The entries in this table are called page table entries (PTE) and store the mapping to physical page number
- Each PTE is **typically** 32 bits (could also be less or more than 32 bits)
 - Assuming 4 KByte page size
 - need at most 20 bits of the PTE for storing the physical page number
 - What about the remaining 12 bits?
 - Can use these remaining bits to store other information e.g. access permissions (whether this physical page stores code/data etc.)

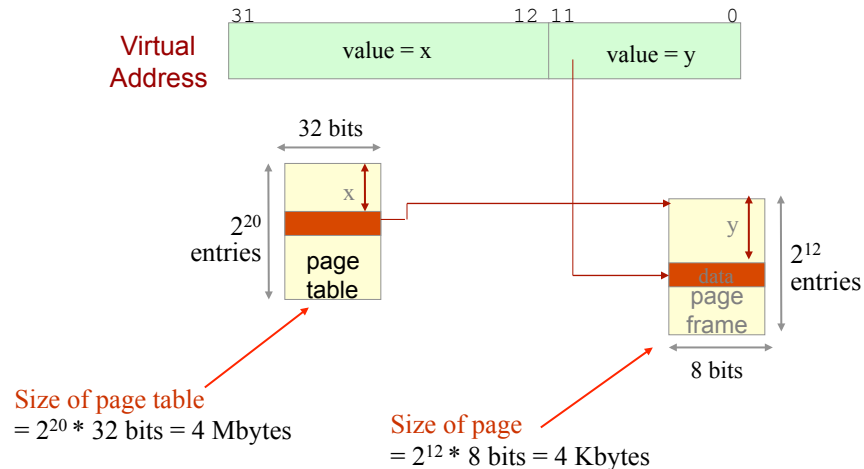
Page Table Entries

- A real **page table entry** (PTE) contains a number of fields
 - Physical page number
 - Valid/Invalid bit
 - Access control bits (e.g., writeable bit)
 - Status bits (e.g., accessed and dirty bits)
 - MMU control bits (e.g., cacheable bits)



Sample page table entry

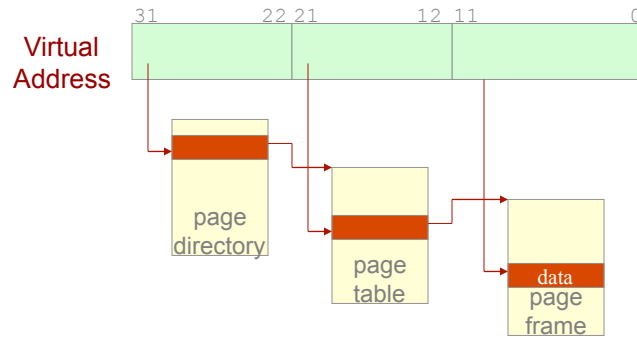
Example: Single-Level Page Table




Single-Level Page Table

- Assumptions
 - 32-bit virtual addresses
 - 4 Kbyte page size = 2^{12} bytes
 - 32-bit address space
- How many virtual page numbers?
 - $2^{32} / 2^{12} = 2^{20} = 1,048,576$ virtual page numbers = number of entries in the page table
- If each page table entry occupies 4 bytes, how much memory is needed to store the page table?
 - $2^{20} \text{ entries} * 4 \text{ bytes} = 2^{22} \text{ bytes} = 4 \text{ Mbytes}$
- How do we reduce the amount of main memory to store the page table?

Two-level Page Tables




Two-Level Page Table Scheme



Source: Silberschatz, Galvin & Gagne

Example: Two-level Page Table




Two-Level Page Table

- Assumptions
 - 2^{10} entries in page directory (= max number of page tables)
 - 2^{10} entries in page table
 - 32 bits allocated for each page directory entry
 - 32 bits allocated for each page table entry
- How much memory is needed?
 - Page table size = $2^{10} \text{ entries} * 32 \text{ bits} = 2^{12} \text{ bytes} = 4 \text{ Kbytes}$
 - Page directory size = $2^{10} \text{ entries} * 32 \text{ bits} = 2^{12} \text{ bytes} = 4 \text{ Kbytes}$

So Where Are We?


- Virtual Memory simplifies memory management by providing each process with same uniform address space



- Protects one process from corrupting data of another process or corrupting its own read only (text) sections
 - The OS can ensure that the VA generated by a process map only to the PA of that process only (unless it explicitly shares code or data)
 - Page Table Entries have access bits that can be used to disallow writes to text or read-only data sections or kernel code/data or shared library

Address translation

- The page table is typically stored in memory
- Steps involved in accessing a page in memory (assuming single-level page table)
 - MMU obtains the address from the CPU & partitions the address into virtual page number VPN and the virtual page offset VPO
 - MMU generates the address of the page table entry (PTE) and requests the contents from memory
 - Memory returns PTE to MMU
 - MMU constructs the physical address and sends it to memory
 - Memory returns the requested data word to CPU

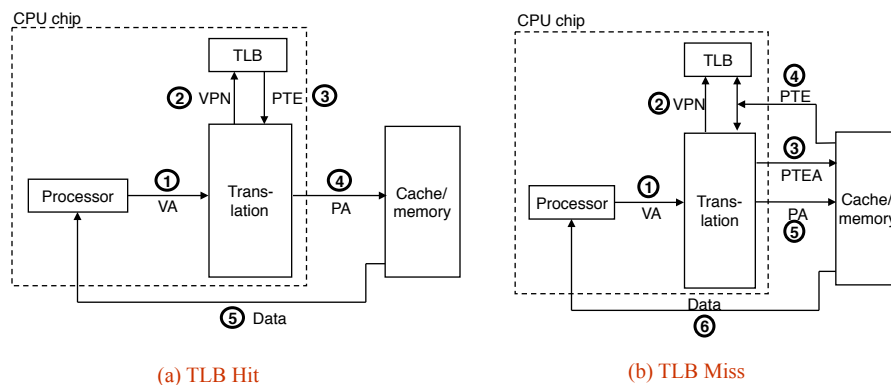


Source: Silberschatz, Galvin & Gagne

Speeding Up Address Translation with TLB

- Where is the page table stored?
- Requires two memory accesses to access any memory location
- Use “Translation Lookaside Buffer” (TLB) to speed up address translation
 - TLB is a small, fast-lookup hardware cache in the MMU
 - TLB contains a few page-table entries
 - When virtual address is generated, it is first presented to TLB
 - If page number is found in TLB (known as **TLB hit**), its frame number is immediately available and is used to access main memory
 - If page number is not found in TLB (known as **TLB miss**), a memory reference to the page table is made


Translation Lookaside Buffer




Source: Bryant & O'Hallaron

Interaction With Caches

- Two possibilities for instruction and data cache (not the TLB) locations
- Logical/Virtual Cache: operates on virtual addresses
 - Advantages/disadvantages?
 - Example: Instruction/Data caches on the XScale processor
- Physical Cache: operates on physical addresses
 - Advantages/disadvantages?
 - Example: Instruction/Data caches on the Intel Pentium processor




Logical cache




Physical cache

Context Switching With Virtual Memory

- Each task in the system has its own page table
 - PTBR needs to be part of the context



Task 1 running



Task 2 running

More On Context Switching

- TLB entries need to be invalidated
 - Since the new process has the same virtual address space (but different physical address space), TLB entries are no longer valid
- Need to clean virtual cache (if any)
 - Cleaning virtual cache may involve more than invalidating the cache item
 - Cache memory on write could follow one of the two policies
 - Writethrough caches: on a cache hit for a write operation, both the cache as well as the main memory is updated
 - Writeback caches: on a cache hit for a write operation, only the cache is updated
 - Memory is updated only when the cache item is evicted
 - For a writeback logical cache, cleaning the cache would involve writing back all the modified cache locations

ARM MMU

- Presents 4 GB virtual address space (why?)
- Memory granularity: 4 options supported
 - 1MB sections
 - Large pages (64 KBytes) - access control within a large page on 16 KBytes
 - Small pages (4 KBytes) - access control within a large page on 1 Kbytes
 - Tiny pages (1KBytes)

31	16	15	12	11	10	3	2	1	0			
Base Physical Address			0000		AP3	AP2	AP1	AP0	C	B	0	1

Large page table entry

ARM MMU (contd.)

- Puts processor in **Abort Mode** when virtual address not mapped or permission check fails
- A page table base register (called the **translation table base register**, in ARM jargon) stores the base address of the page table
 - Useful for context switching of processes
- XScale processor (used on the XBoard) has
 - 32 entry instruction and data TLB
 - 32 Kbyte instruction and data cache

Summary of Lecture

- Virtual Memory (VM)
 - What is VM?
 - How does VM work?
 - Virtual address to physical address mappings
 - Page faults
 - VM Schemes
 - page tables
 - virtual pages and physical page frames
 - page table entries
 - TLBs
 - Interaction with caches
- **References**
 - R. E. Bryant & D. R. O'Hallaron, "*Computer Systems: A Programmer's Perspective*" Chapter 10
 - Silberschatz, Galvin and Gagne, "*Operating System Concept*" Chapters 8 & 9
 - A. S. Tannenbaum, "*Modern Operating Systems*", Chapter 4